

# Kubernetes Network Policy and Istio Service Mesh - Student Guide

## Overview

This lesson demonstrates two important Kubernetes networking concepts:

1. **Network Policies** - How to control network traffic between pods
2. **Istio Service Mesh** - Advanced traffic management with sidecar pattern

## Learning Objectives

By the end of this lesson, you will be able to:

- Configure Kubernetes Network Policies to control pod-to-pod communication
  - Set up Istio service mesh in a Kubernetes cluster
  - Understand the sidecar pattern and its benefits
  - Deploy and access applications through Istio ingress gateway
- 

## Part 1: Network Policy Demo

### Objective

Block frontend pods from connecting to MySQL database while allowing backend pods to maintain connectivity.

### Prerequisites

- Minikube installed
- kubectl configured
- Basic understanding of Kubernetes pods and services

### Step 1: Set up Minikube with Calico CNI

First, clean up any existing Minikube cluster:

```
bash
minikube delete
```

Start Minikube with Calico CNI (required for Network Policy support):

```
bash

minikube start --network-plugin=cni --cni=calico
```

**Why Calico?** Network Policies require a CNI plugin that supports them. Calico is one of the most popular choices.

Verify Calico is running:

```
bash

kubectl get pods -l k8s-app=calico-node -n kube-system
```

## Step 2: Deploy the Application

Apply the demo application manifest:

```
bash

kubectl apply -f network-policy-demo.yaml
```

This creates:

- **Frontend pod** - Simulates a web frontend
- **Backend pod** - Simulates an API backend
- **MySQL pod** - Database server
- **Services** - ClusterIP services for each component

Verify all pods are running:

```
bash

kubectl get pods
```

Expected output:

```
NAME      READY   STATUS    RESTARTS   AGE
backend    1/1     Running   0           56s
frontend  1/1     Running   0           56s
mysql      1/1     Running   0           56s
```

## Step 3: Test Initial Connectivity

Check that frontend can connect to database (this should work initially):

```
bash  
kubectl exec -it frontend -- bash
```

Inside the frontend pod:

```
bash  
apt update && apt install telnet -y  
telnet db 3306
```

You should see:

```
Trying 10.105.255.201...  
Connected to db.
```

Exit the telnet session and pod:

```
bash  
exit  
exit
```

## Step 4: Apply Network Policy

Now apply the network policy to block frontend access:

```
bash  
kubectl apply -f db-netpol.yaml
```

## Step 5: Verify the Policy Works

Test that frontend is now blocked:

```
bash  
kubectl exec -it frontend -- bash  
telnet db 3306
```

This should hang (connection blocked). Use **Ctrl+C** to exit.

Test that backend still works:

```
bash
kubectl exec -it backend -- bash
apt update && apt install telnet -y
telnet db 3306
```

This should still connect successfully.

## Understanding the Network Policy

The `db-netpol.yaml` file contains rules that:

- Apply to pods with specific labels (the database)
  - Allow ingress traffic only from pods with certain labels (backend)
  - Block all other traffic by default
- 

## Part 2: Istio Service Mesh Demo

### Objective

Deploy Istio service mesh and demonstrate the sidecar pattern with the Bookinfo application.

### Step 1: Prepare Minikube for Istio

Start Minikube with sufficient memory:

```
bash
minikube start --memory=4096
```

**Why more memory?** Istio components require additional resources to run properly.

### Step 2: Download and Install Istio

Download Istio:

```
bash
curl -L https://istio.io/downloadIstio | sh -
```

Navigate to the Istio directory:

```
bash
```

```
cd istio-1.26.2/bin
```

Install Istio with demo profile:

```
bash
```

```
./istioctl install --set profile=demo -y
```

The demo profile includes:

- Istio core components
- Ingress and egress gateways
- Additional features for learning

### Step 3: Configure Automatic Sidecar Injection

Enable automatic sidecar injection for the default namespace:

```
bash
```

```
kubectl label namespace default istio-injection=enabled
```

**What is sidecar injection?** Istio automatically injects a proxy container (Envoy) alongside each application container to handle network traffic.

### Step 4: Deploy the Bookinfo Application

Deploy the sample Bookinfo application:

```
bash
```

```
kubectl apply -f samples/bookinfo/platform/kube/bookinfo.yaml
```

This creates a microservices application with:

- **Product page** - Main application frontend
- **Details** - Book details service
- **Reviews** - Book reviews service (3 versions)
- **Ratings** - Star ratings service

## Step 5: Verify Sidecar Injection

Check that pods have 2/2 containers (app + sidecar):

```
bash
kubectl get pods
```

Expected output:

NAME	READY	STATUS	RESTARTS	AGE
details-v1-766844796b-jkbrg	2/2	Running	0	77s
productpage-v1-54bb874995-q4c6g	2/2	Running	0	77s
ratings-v1-5dc79b6bcd-55fkn	2/2	Running	0	77s
reviews-v1-598b896c9d-85n8t	2/2	Running	0	77s
reviews-v2-556d6457d-l7ktl	2/2	Running	0	77s
reviews-v3-564544b4d6-9lr8t	2/2	Running	0	77s

**Notice:** Each pod shows `2/2` ready, indicating the application container plus the Istio sidecar.

## Step 6: Configure Ingress Gateway

Apply the gateway configuration:

```
bash
kubectl apply -f samples/bookinfo/networking/bookinfo-gateway.yaml
```

This creates:

- **Gateway** - Configures the ingress gateway
- **VirtualService** - Routes traffic to the product page

## Step 7: Access the Application

Get the ingress gateway URL:

```
bash
minikube service istio-ingressgateway -n istio-system --url
```

This returns multiple URLs. Use the second one (typically on port 80) to access:

<http://192.168.49.2:31171/productpage>

Open this URL in your browser to see the Bookinfo application.

---

## Key Concepts Explained

### Network Policies

- **Purpose:** Control traffic flow between pods at the network level
- **Default behavior:** Without policies, all pods can communicate
- **Label-based:** Use selectors to define which pods the policy applies to
- **Ingress/Egress:** Control incoming and outgoing traffic separately

### Istio Service Mesh

- **Sidecar Pattern:** Each pod gets a proxy container that handles network traffic
- **Traffic Management:** Control routing, load balancing, and failover
- **Security:** Mutual TLS, authentication, and authorization
- **Observability:** Metrics, logging, and tracing out of the box

### Benefits of Service Mesh

- **Decoupled:** Network logic separated from application code
  - **Consistent:** Same features across all services
  - **Secure:** Automatic encryption and authentication
  - **Observable:** Built-in monitoring and tracing
- 

## Troubleshooting Tips

### Network Policy Issues

- Ensure you're using a CNI that supports Network Policies (like Calico)
- Check pod labels match the policy selectors
- Remember: policies are namespace-scoped

### Istio Issues

- Verify sufficient memory allocation to Minikube
- Check that sidecar injection is enabled for the namespace

- Ensure all pods show 2/2 ready status
- Use `istioctl analyze` to diagnose configuration issues

## Common Commands

```
bash
```

```
# Check network policies
```

```
kubectl get networkpolicies
```

```
# Describe a network policy
```

```
kubectl describe networkpolicy <policy-name>
```

```
# Check Istio configuration
```

```
istioctl analyze
```

```
# View Istio proxy configuration
```

```
istioctl proxy-config cluster <pod-name>
```

---

## Exercise Questions

1. What happens if you remove the network policy? Test it.
2. How would you modify the policy to allow frontend access on a different port?
3. What's the difference between pods with and without Istio sidecars?
4. How can you verify that traffic is being encrypted by Istio?

---

## Additional Resources

- [Kubernetes Network Policy Documentation](#)
- [Istio Documentation](#)
- [Calico Network Policy Guide](#)
- [Envoy Proxy Documentation](#)